



# From Zero to Hero with Apache Kudu

**BOSS 2019**

Andrew Wong

Follow along with the walkthroughs!

<https://bit.ly/2ZsR6m6>

---

# Who am I?

Andrew Wong

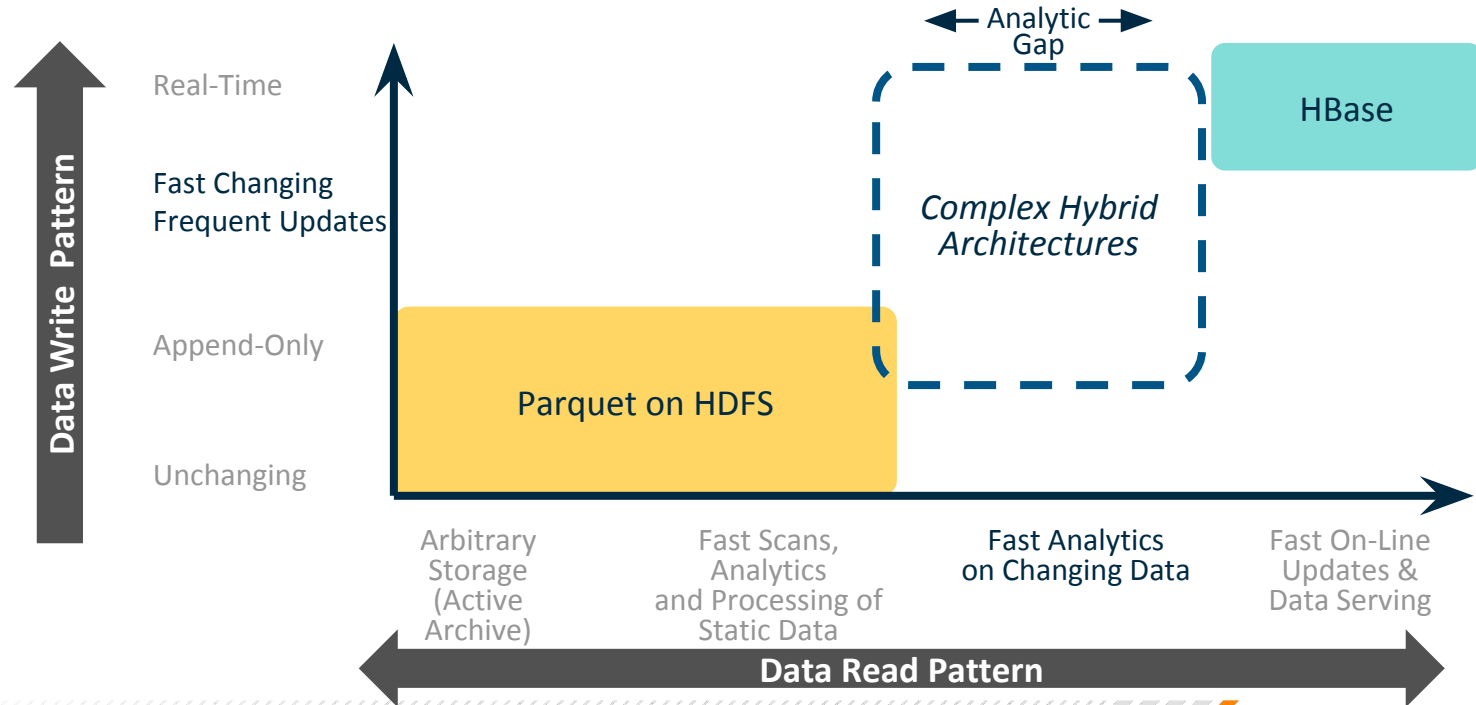
- Software engineer at Cloudera ([awong@cloudera.com](mailto:awong@cloudera.com))
- Apache Kudu PMC Member

Some Kudu things I have worked on:

- Scan optimizations
- Disk failure mitigation
- Integration with Hive Metastore (external catalog)
- Fine-grained authorization

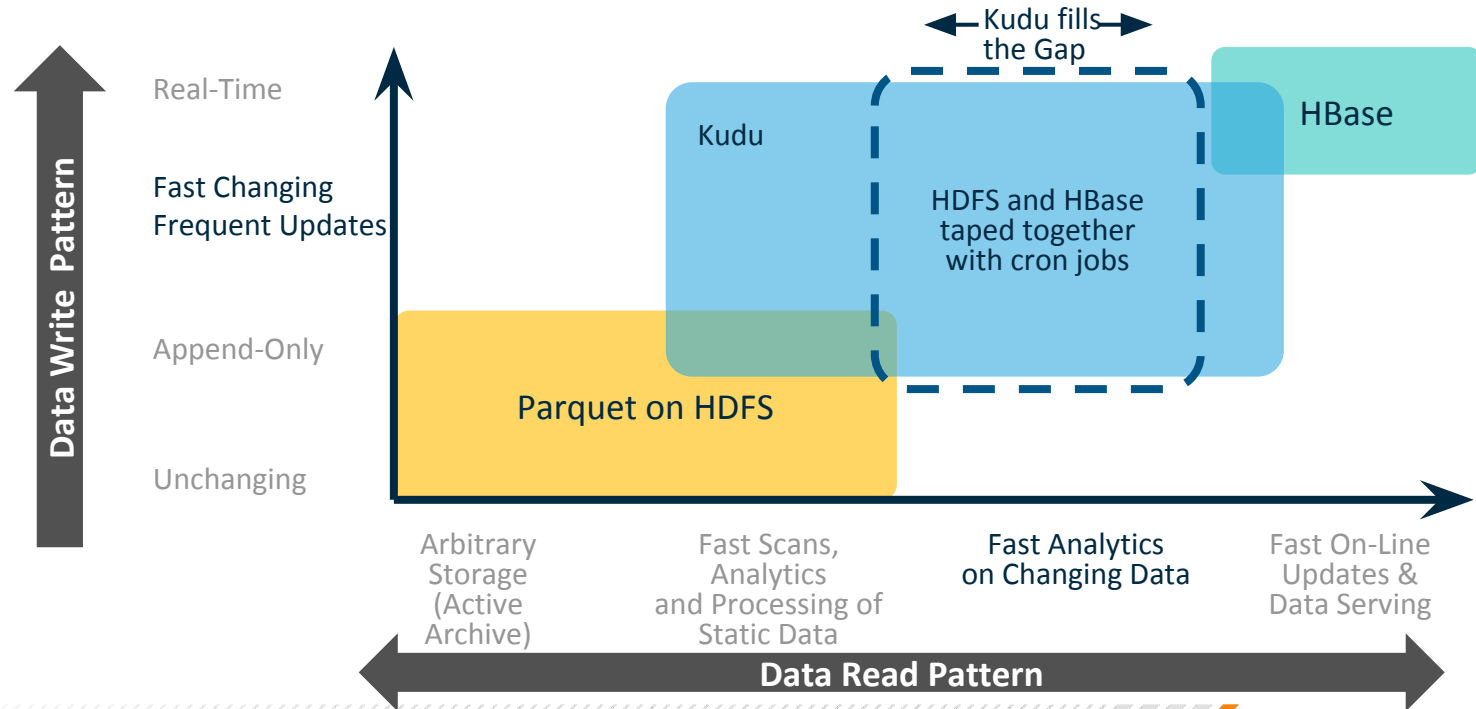
# Traditional big data storage leaves a gap

Use cases fall between HDFS and HBase were difficult to manage

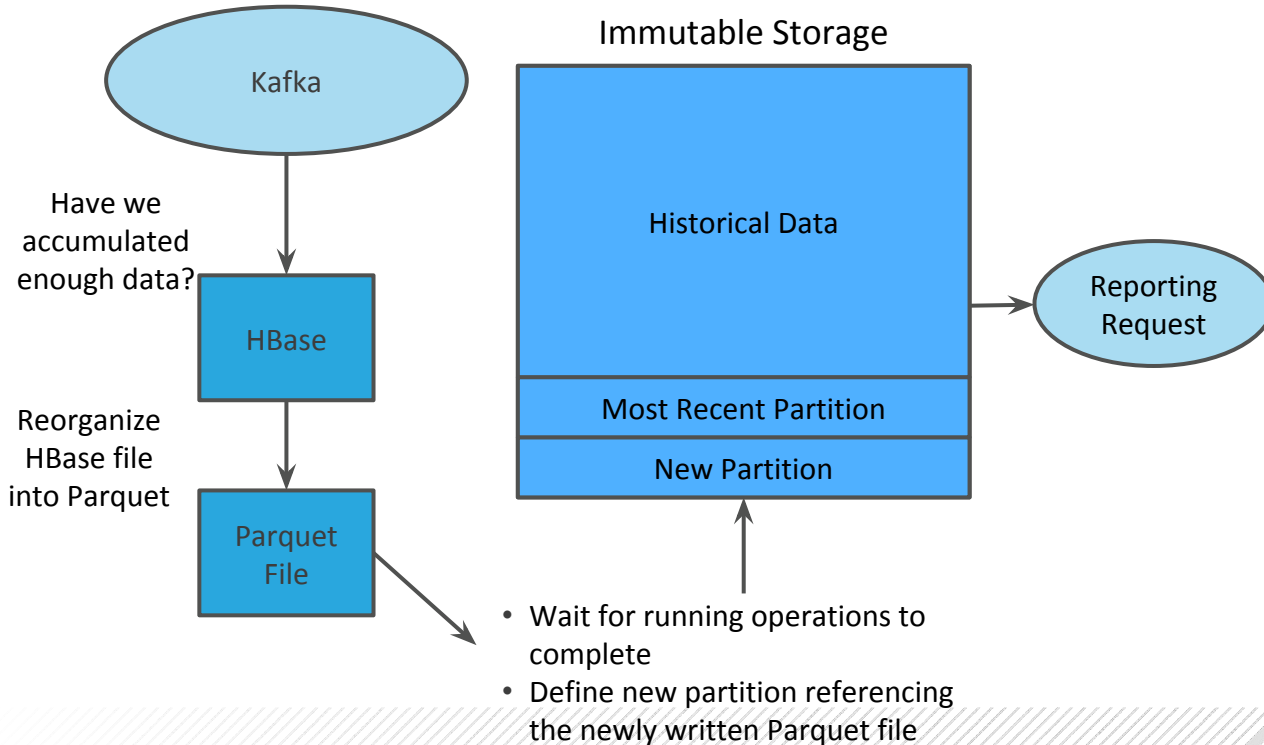


# Traditional big data storage leaves a gap

Use cases fall between HDFS and HBase were difficult to manage



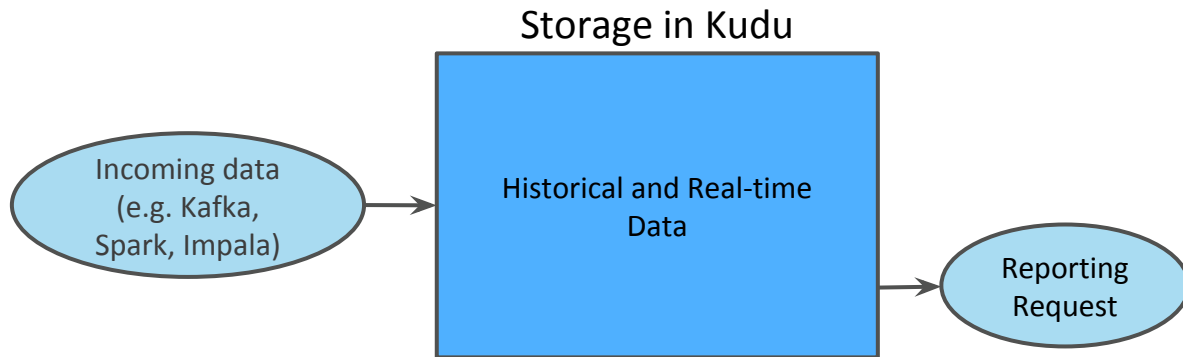
# “Traditional” real-time analytics



## Considerations:

- How do I handle failure during this process?
- How often do I reorganize data streaming in into a format appropriate for reporting?
- When reporting, how do I see data that has not yet been reorganized?
- How do I ensure that important jobs aren't interrupted by maintenance?

# Real-time analytics with Kudu



## Improvements:

- Much simpler architecture
- Significantly easier to handle late arrivals of data
- New data available immediately for analytics or operations

---

# What is Apache Kudu?

Mutable data storage engine, designed for analytics on real-time data

MUTABLE

STORAGE

ANALYTICS

REAL-TIME DATA

# What is Apache Kudu?

Mutable data storage engine, designed for analytics on real-time data

MUTABLE

STORAGE

ANALYTICS

REAL-TIME DATA





# What is Apache Kudu?

Mutable data storage engine, designed for analytics on real-time data

MUTABLE

STORAGE

ANALYTICS

REAL-TIME DATA



# What is Apache Kudu?

Mutable data storage engine, designed for analytics on real-time data

MUTABLE

STORAGE

ANALYTICS

REAL-TIME DATA



# What is Apache Kudu?

Mutable data storage engine, designed for analytics on real-time data

MUTABLE

STORAGE

ANALYTICS

REAL-TIME DATA



---

# “Big deal... Aren’t there a ton of big data systems out there?”

Yes there are.

- Open source
  - Parquet on HDFS, object storage
  - HBase/Cassandra
  - TiDB with TiFlash
  - etc.
- Proprietary
  - Vertica, Teradata, SAP
  - Spanner
  - Redshift

---

# Why else did we build Kudu?

## Changing hardware landscape

- Spinning disks --> solid state storage
  - **NAND flash:** up to 450k read, 250k write IOPS, ~2GB/s read and ~1.5GB/s write throughput, at under \$1/GB
  - **PMEM:** order of magnitude faster than NAND, cheaper than RAM
- **RAM** is getting cheaper and more abundant
  - 128 --> 256 --> 512 GB over the last few years

**Takeaway:** The next bottleneck is CPU, and current storage systems weren't designed with CPU efficiency in mind.

---

# Scalable and fast tabular storage

## Scalable

- Production clusters with hundreds of nodes
- On the order of hundreds of TBs to low PBs

## Fast


- Written primarily in C++
- Millions of write operations per second across cluster
- Multiple GB/s read throughput per node

## Tabular

- Strict schema, finite column count, no BLOBs

And we've got some friends!

Apache and CNCF



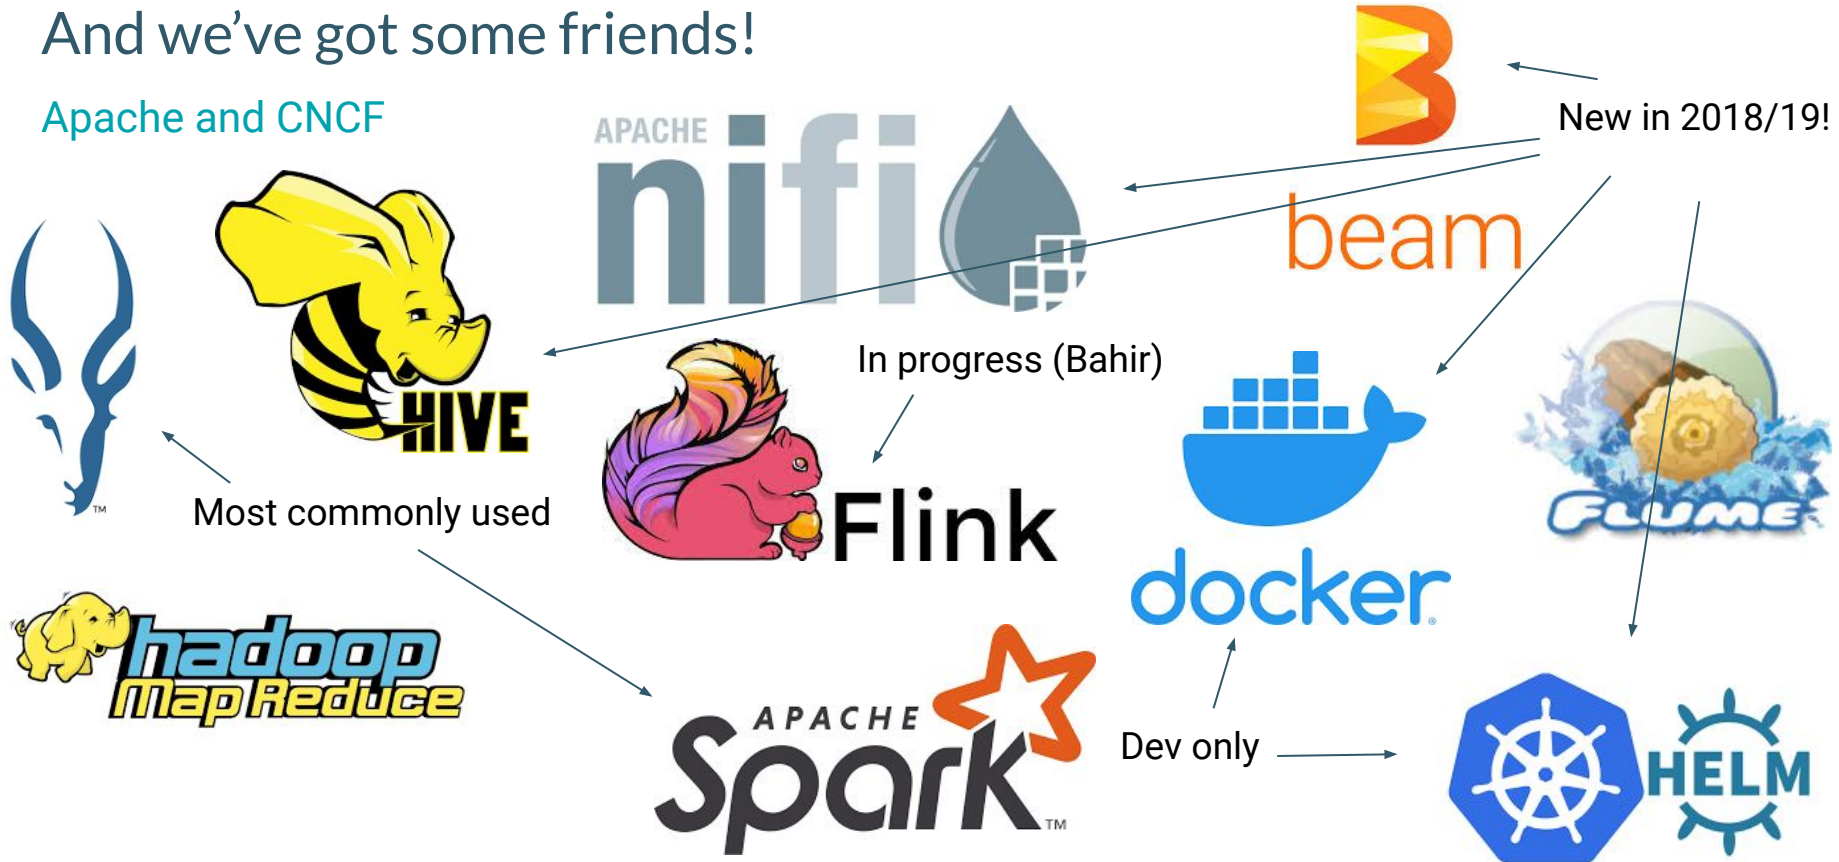
beam





# And we've got some friends!

Apache and CNCF





# The rest of the Ecosystem isn't that scary

Somewhat clear lines between roles

Ingest:

- NiFi, Spark Streaming, Flink, Kafka, Flume



Storage:

- **Kudu**
- Previously HBase, HDFS, and a lot of cron jobs

Querying:

- Impala, Hive, Spark SQL



---

# From Zero to Hero with Apache Kudu

## Agenda

Start with the basics:

- Introduction to Kudu's data model
- Distributed architecture

Jump into more complex things:

- Schema design exercises
- Partitioning in Kudu

Deploy a small “big data” pipeline:

- Spark, Nifi

---

# From Zero to Hero with Apache Kudu

If you can

- Pull from GitHub:
  - `apache/kudu:master`
- Pull from DockerHub:
  - `apache/kudu:latest`
- `brew install apache-spark`

---

# A primer on Kudu concepts

## What is “data” to Kudu?

**Table:** prominent abstraction for users -- a set of uniquely identifiable rows

**Schema:** describes the columns and ordering of the rows of a table

- Subset of columns defined as “primary key”

**Partition schema:** describes the partitions of a table

- Subset of primary key defined as “partition key”

**Tablet:** partition of a table; the logical unit of replication and parallelization

- Underlying data stored in sorted order by primary key
- Tablet replicas are dispersed among servers in a cluster

---

# What makes a tablet?

Very much LSM-inspired

**Write-ahead log:** as writes come in, they are written in fast, row-oriented storage

**Mem-Rowset:** as writes are written to the WAL, they are also applied to an in-memory, row-oriented B-Tree

**Periodic flushes:** as the mem-rowset grows large, there are periodic flushes to free up memory and transition the mem-rowset to a disk-rowset

**Many Disk-Rowsets:** represented as an interval tree for lookups

---

# What makes a tablet?

Wait what is “LSM”? You said we’re starting from zero!

- LSM: Log Structured Merge (Cassandra, HBase, etc)
  - Inserts and updates all go to an in-memory map (MemStore) and later flush to on-disk files (HFile/SSTable)
  - Reads perform an on-the-fly merge of all on-disk HFiles
- Kudu
  - Shares some traits (memstores, compactions)
  - More complex.
  - **Slower writes** in exchange for **faster reads** (especially scans)

---

# What makes a tablet?

And more!

**B-Tree index** per rowset to enable search on a primary key

**Bloom filter** per rowset to enable lookups for deduplication of rows

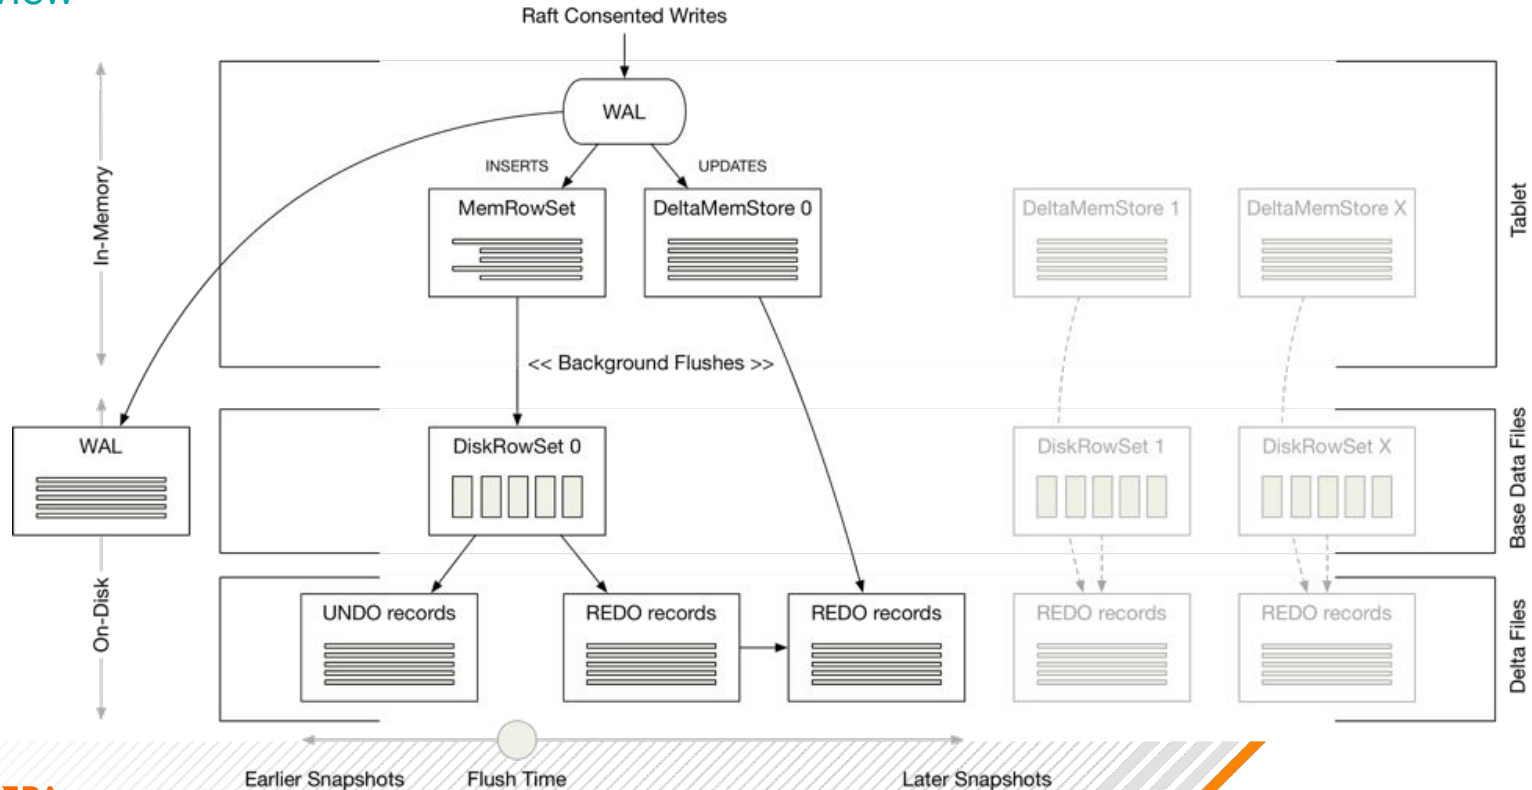
**CFiles** the actual columnar data, in primary key order

Deltas (updates and deletes)

- **In-memory delta store** per rowset
- **On-disk delta stores** per rowset
- New updates are represented as REDO records, periodically flushed to UNDO records to avoid REDO traversal

# What makes a tablet?

## Overview





---

# Maintenance Manager

Always be flushing and compacting!

- The maintenance manager handles background tasks
  - Flushing the MRS or DMS to disk
  - Compactions
  - WAL GC
- A maintenance manager thread decides what task to perform using a cost-based optimization model
  - Prefers flushing when the server is under memory pressure

---

# Distributed architecture

## Main roles in Kudu

**Master:** a server that hosts Kudu system metadata and catalog information

**Tablet server:** a server that hosts tablet replicas

**Client:** an application that inserts to or reads from Kudu

- Tooling
- C++ application
- Java
- Python
- Spark
- Impala
- ...

---

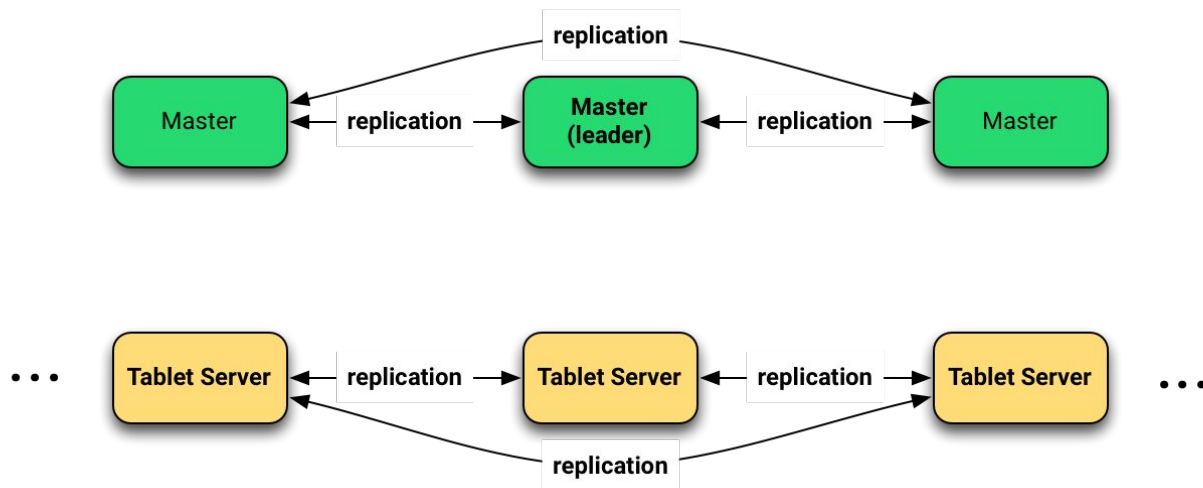
# Let's walk through deploying Kudu!

<https://bit.ly/2ZsR6m6>

- Explore a single node deployment of Kudu

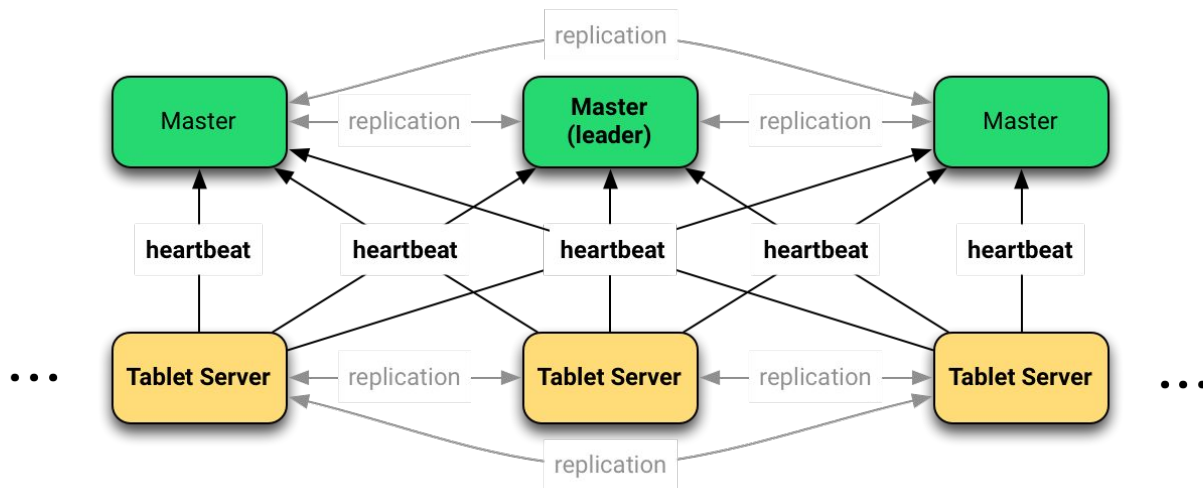
# Distributed architecture

Replication, replication, replication!



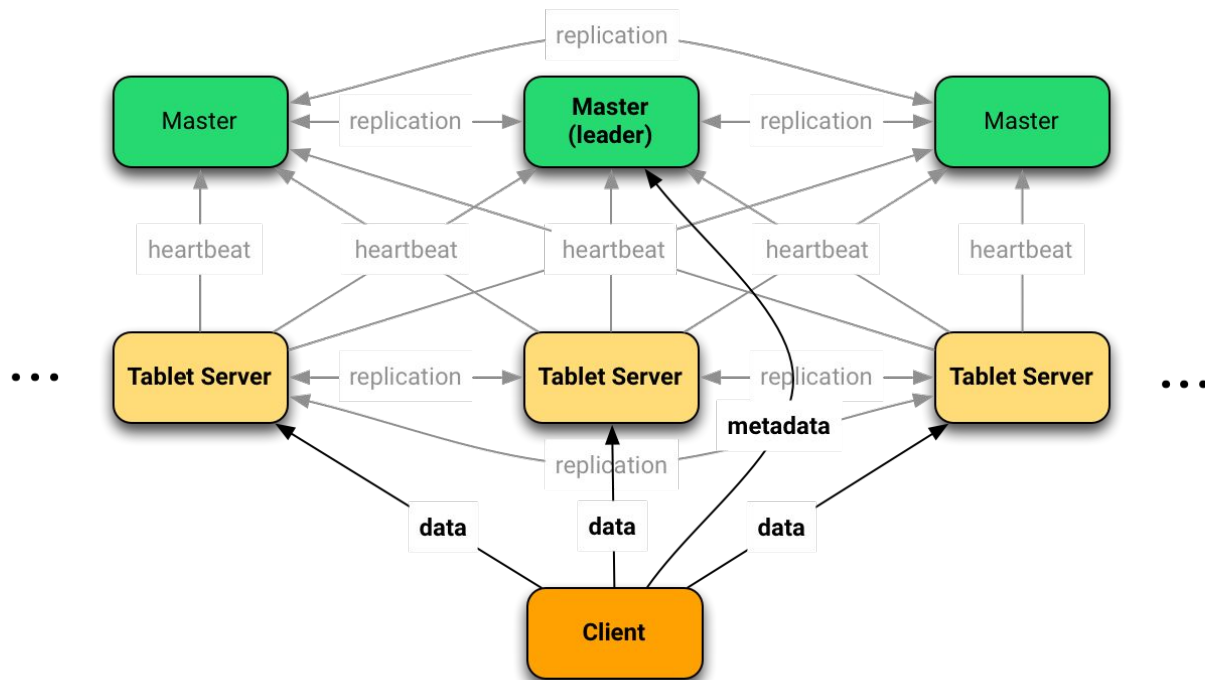
# Distributed architecture

Replication, replication, replication!



# Distributed architecture

Replication, replication, replication!



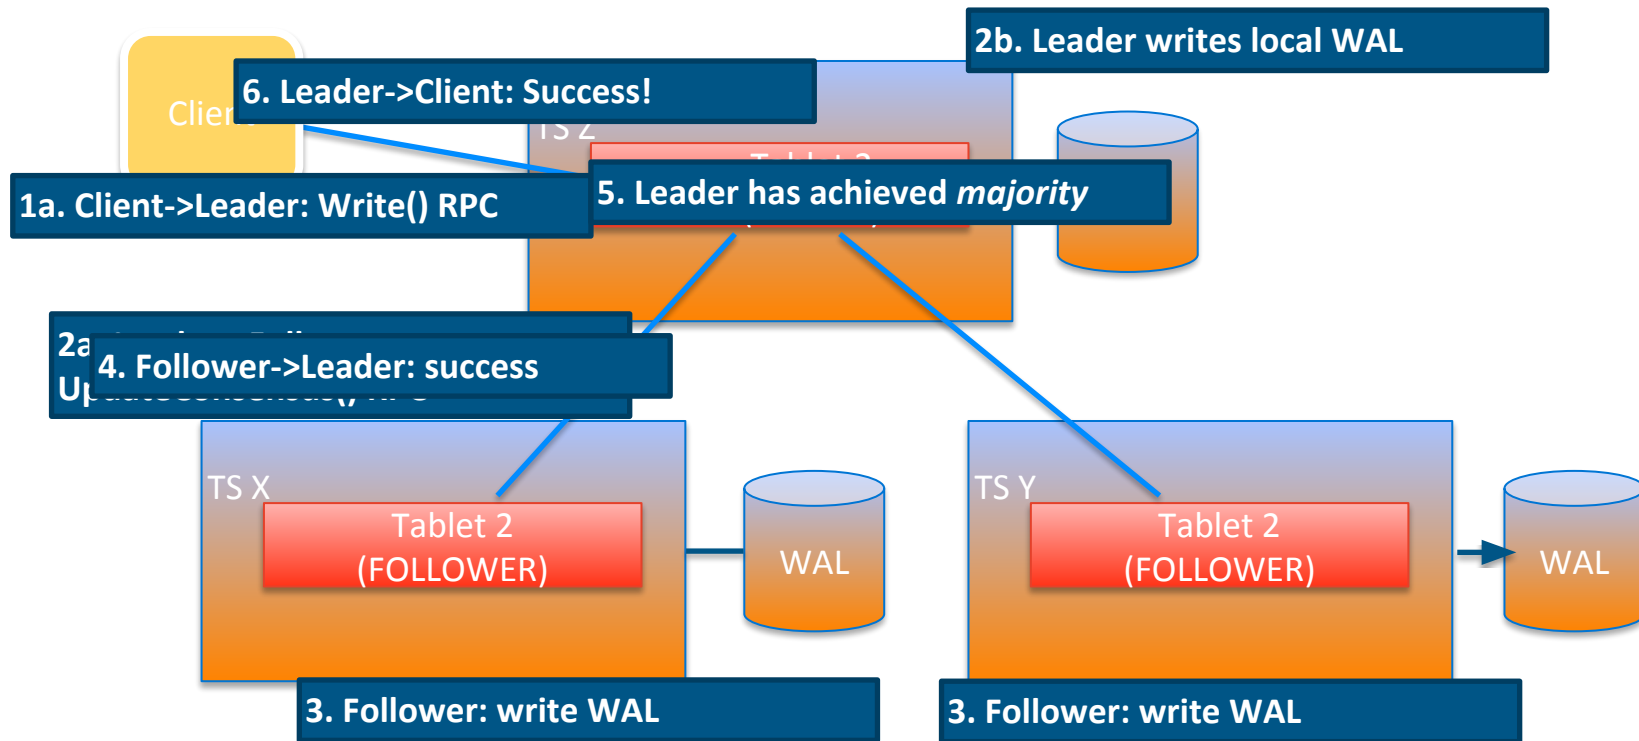
# Raft consensus

## Consensus protocol to replicate data

- Leaders constantly heartbeating to followers
- Leader elections triggered when a follower doesn't hear back from leader
- Write operations go to leader first and are replicated to followers



# Raft consensus





# Dealing with node failures

Wait long enough, and Kudu heals itself!

- If a Raft leader notices that it hasn't heard from one of its replicas in a while, it will try to create a new replica
- Physical copy of leader sent to new replica
- Old replica is removed from Raft configuration



---

## Master details

- All table and tablet metadata is stored in a single tablet
  - Special tablet ID: 000000000000000000000000000000000000
- Like any other tablet, the metadata tablet is replicated via Raft
- Master is effectively a tablet server for this one tablet
- Failure handling
  - Transient failures are handled transparently
  - Permanent failures require operator intervention

---

# Balance and Skew

Tablet replica placement controlled by the master based on its view of the world

- Placement uses “Power of Two Choices” algorithm to even out replica count across cluster
  - Select two tablet servers, pick the one with fewer tablet replicas
- Rebalancer tool can be run to redistribute replicas in case of skew

---

# Cluster health

`kudu cluster ksck <masters>`

When operating Kudu, this is your lifeline!

Tells you:

- What servers exist
- What tables and tablets exist
- What tablet are under-replicated and recovering
- What the skew on the cluster is like
- What non-default flags are set

---

# Client abstractions

## Sessions:

- In-memory write buffers, eventually get “flushed” to the servers

## Scanners:

- Fetches rows one batch at a time

---

## Client abstractions (writes)

```
KuduTable table = client.openTable("metrics");
KuduSession session = client.newSession();
Insert ins = table.newInsert();
ins.getRow().addString("host", "foo.example.com");
ins.getRow().addString("metric", "load-avg.1sec");
ins.getRow().addDouble("value", 0.05);
session.apply(ins);
session.flush();
```

---

## Client abstractions (scans)

```
KuduScanner scanner = client.newScannerBuilder(table)
    .setProjectedColumnNames(Lists.of("value"))
    .build();
while (scanner.hasMoreRows()) {
    RowResultIterator batch = scanner.nextRows();
    while (batch.hasNext()) {
        RowResult result = batch.next();
        System.out.println(result.getDouble("value"));
    }
}
```

---

## Client abstractions (scans, but with predicates!)

```
KuduScanner scanner = client.newScannerBuilder(table)
    .addPredicate(KuduPredicate.newComparisonPredicate(
        table.getSchema().getColumn("timestamp"),
        ComparisonOp.GREATER,
        System.currentTimeMillis() / 1000 + 60))
    .build();
```

Note: Kudu can evaluate simple predicates, but no aggregations, complex expressions, UDFs, etc.



---

# Client cluster interaction

## Master

- Tablet metadata fetched from the master, including locations and partitioning info
- Tablets are “pruned” via partitioning info so only the appropriate tablets are scanned

## Tablet servers

- Requests sent to tablet servers
- Rejected if appropriate

---

# Client consistency models

Choose the one which fits your workload!

- READ\_LATEST (default mode)
  - Read committed state immediately
- READ\_AT\_SNAPSHOT
  - Consistent and repeatable
  - This allows strict-serializable semantics for reads and writes
- READ\_YOUR\_WRITES (Kudu 1.7 and up)
  - Ensures all previously read and written values are read
  - Not repeatable

Meta Cache

T1: ...

T2: ...

T3: ...

Client

Hey Master! Where is the row for  
'tlipton' in table "T"?

It's part of tablet 2, which is on servers {Z,Y,X}.  
BTW, here's info on other tablets you might  
care about: T1, T2, T3, ...

UPDATE tlipton  
SET col=foo

Master tablet

Master  
Server A



Master  
Server B



Master  
Server C



Tablet 1

Tablet 2

...

Tablet n



Tablet  
Server W



Tablet  
Server X



Tablet  
Server Y



Tablet  
Server Z

---

# Continuing the walkthrough!

<https://bit.ly/2ZsR6m6>

- Explore a multinode deployment

# Time Series

---

To demonstrate advanced table partitioning techniques, we are going to think through a table for time series storage of machine metrics.

Series → Time → Value

# Time Series

Series	Time	Value
us-east.appserver01.loadavg.1min	2016-05-09T15:14:30Z	0.44
us-east.appserver01.loadavg.1min	2016-05-09T15:14:40Z	0.53
us-west.dbserver03.rss	2016-05-09T15:14:30Z	1572864
us-west.dbserver03.rss	2016-05-09T15:15:00Z	2097152

---

## Time Series — Design Criteria

- Insert Performance (throughput & latency)
- Read Performance (throughput & latency)
  - What kind of queries are you doing?
    - Look at all metrics at a specific time
    - Look at one metric across a long span of time

# Time Series — Common Patterns

- Datapoints are inserted **in time order** across all series
- Reads specify a **series** and a **time range**, containing hundreds to many thousands of datapoints

```
SELECT time, value FROM timeseries  
WHERE series = "us-west.dbserver03.rss"  
      AND time >= 2016-05-08T00:00:00;
```



---

## Reminder: Partitioning vs Indexing

- **Partitioning:** how datapoints are distributed among partitions
  - Kudu: tablet
  - HBase: region
  - Cassandra: VNode
- **Indexing:** how data within a single partition is sorted

# Reminder: Partitioning vs Indexing

- **Partitioning:** how datapoints are distributed among partitions
  - Kudu: tablet
  - HBase: region
  - Cassandra: VNode
- **Indexing:** how data within a single partition is sorted

```
(us-east.appserver01.loadavg, 2016-05-09T15:14:00Z)
(us-east.appserver01.loadavg, 2016-05-09T15:15:00Z)
(us-west.dbserver03.rss, 2016-05-09T15:14:30Z)
(us-west.dbserver03.rss, 2016-05-09T15:14:30Z)
```

(series, time)

```
(2016-05-09T15:14:00Z, us-east.appserver01.loadavg)
(2016-05-09T15:14:30Z, us-west.dbserver03.rss)
(2016-05-09T15:15:00Z, us-east.appserver01.loadavg)
(2016-05-09T15:14:30Z, us-west.dbserver03.rss)
```

(time, series)

# Reminder: Partitioning vs Indexing

- **Partitioning:** how datapoints are distributed among partitions
  - Kudu: tablet
  - HBase: region
  - Cassandra: VNode
- **Indexing:** how data within a single partition is sorted

```
(us-east.appserver01.loadavg, 2016-05-09T15:14:00Z)  
(us-east.appserver01.loadavg, 2016-05-09T15:15:00Z)  
(us-west.dbserver03.rss, 2016-05-09T15:14:30Z)  
(us-west.dbserver03.rss, 2016-05-09T15:14:30Z)
```

(series, time)

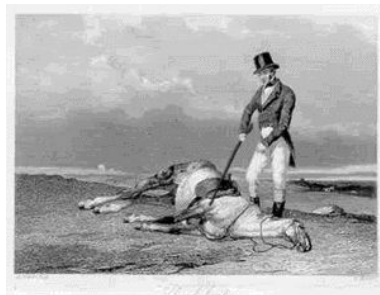
SELECT \* WHERE series = 'us-east.appserver01.loadavg';

```
(2016-05-09T15:14:00Z, us-east.appserver01.loadavg)  
(2016-05-09T15:14:30Z, us-west.dbserver03.rss)  
(2016-05-09T15:15:00Z, us-east.appserver01.loadavg)  
(2016-05-09T15:14:30Z, us-west.dbserver03.rss)
```

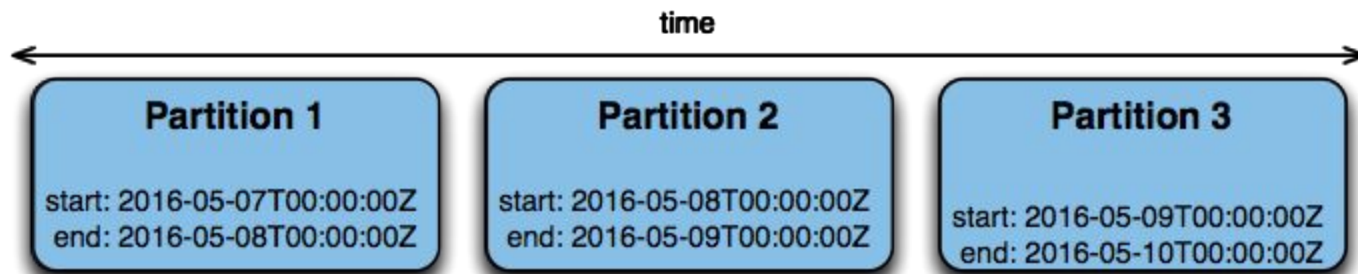
(time, series)

# Partitioning

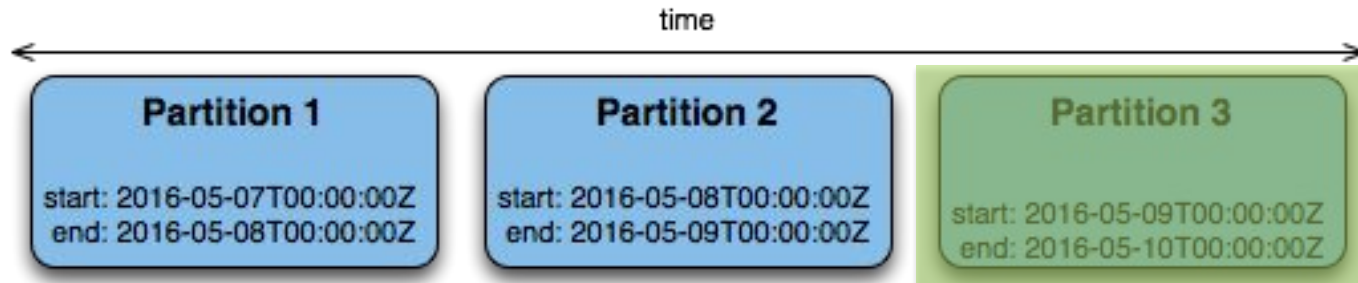
- Kudu has flexible policies for distributing data among partitions
  - Hash partitioning is built in, and can be combined with range partitioning
- **Goal:** acceptable distribution of data across tablets at any given time
- **Goal:** allow expected scans to prune tablets
- Indexing is independent of partitioning!!!



## Partitioning – By Time Range

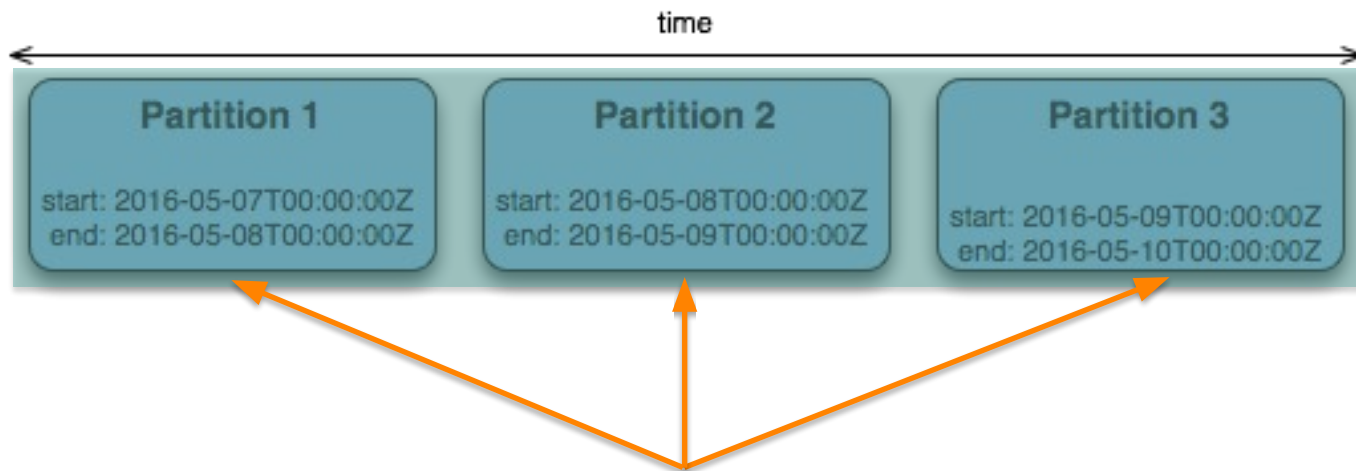


## Partitioning – By Time Range (inserts)



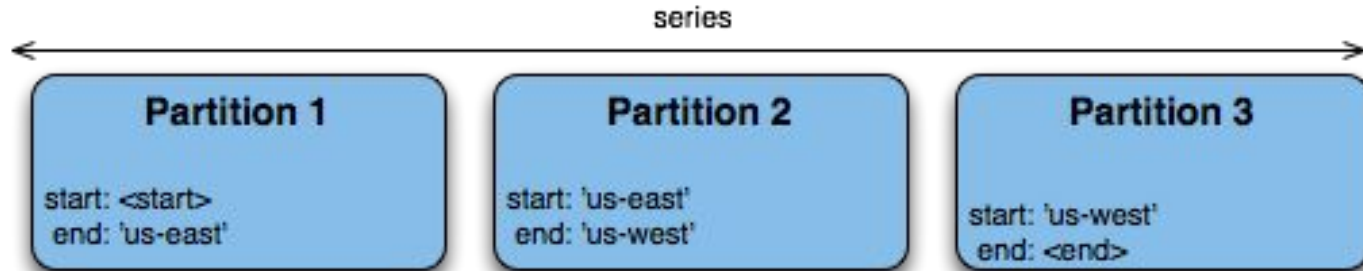
All Inserts go to Latest Partition

## Partitioning – By Time Range (scans)



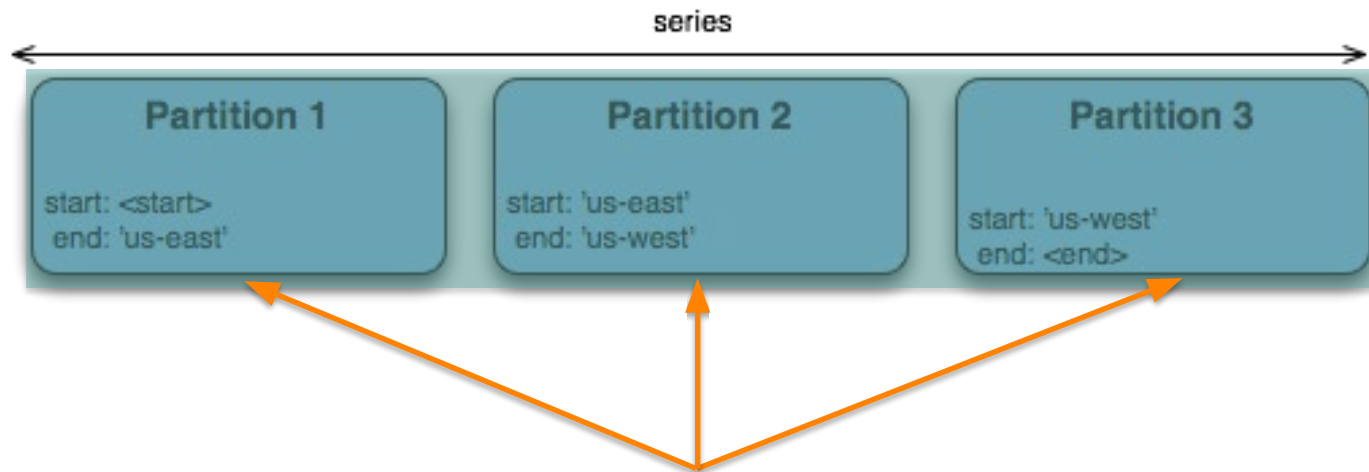
Big scans (across large time intervals)  
can be parallelized across many partitions

## Partitioning — By Series Range



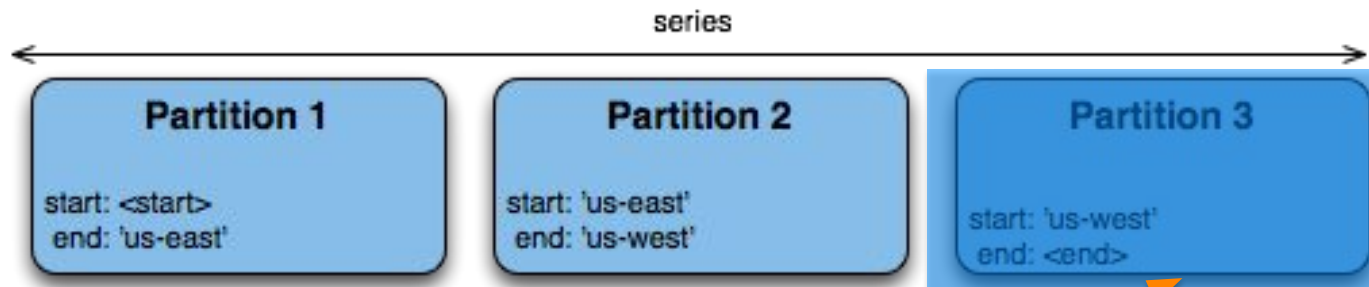


## Partitioning — By Series Range (inserts)



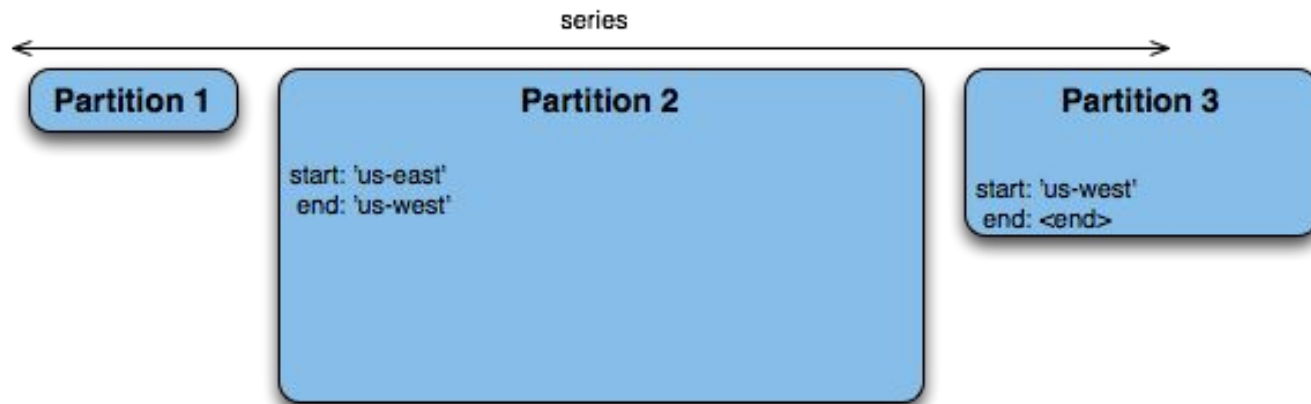
Inserts are spread among all partitions

## Partitioning — By Series Range (scans)



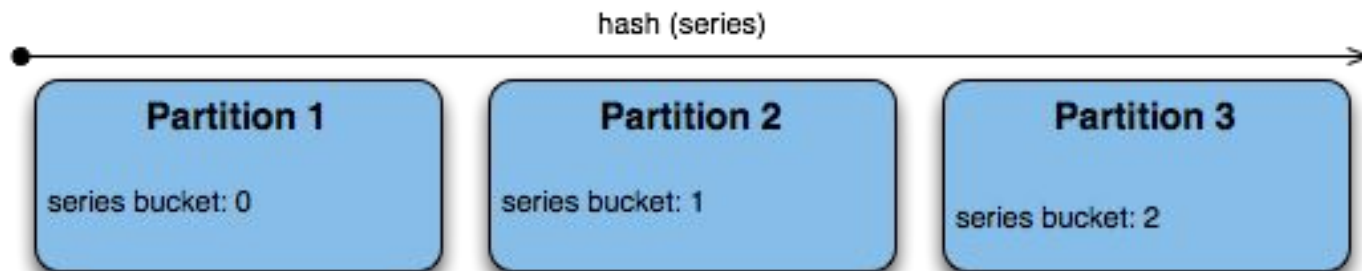
Scans are over a single partition

## Partitioning — By Series Range

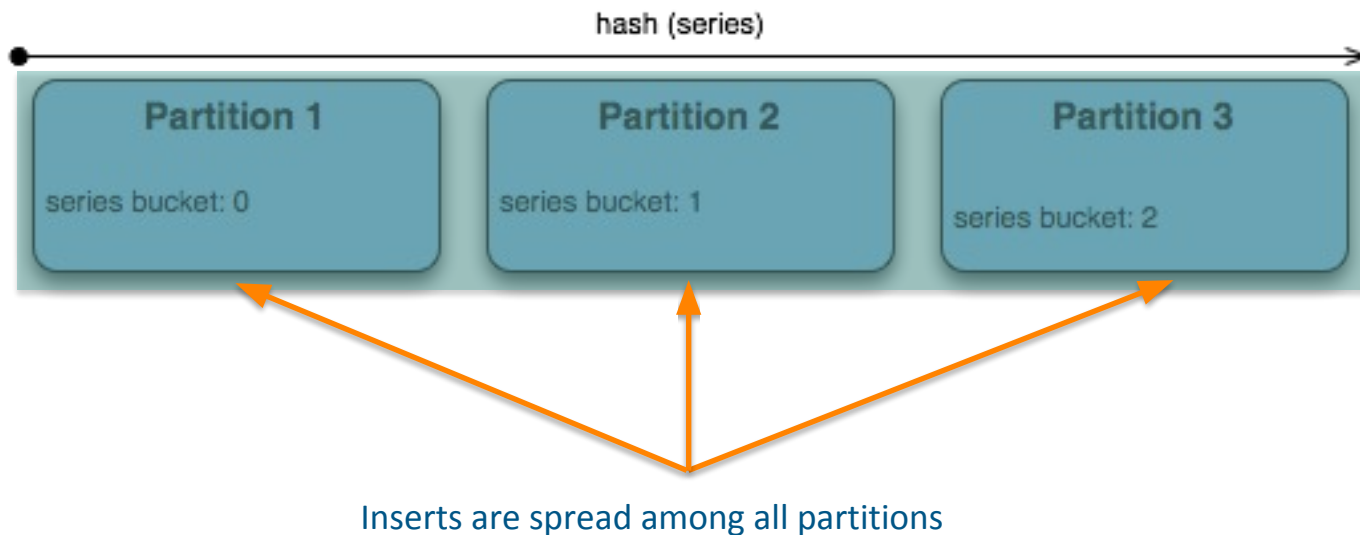


Partitions can become unbalanced,  
resulting in hot spotting

## Partitioning – By Series Hash



## Partitioning – By Series Hash (inserts)

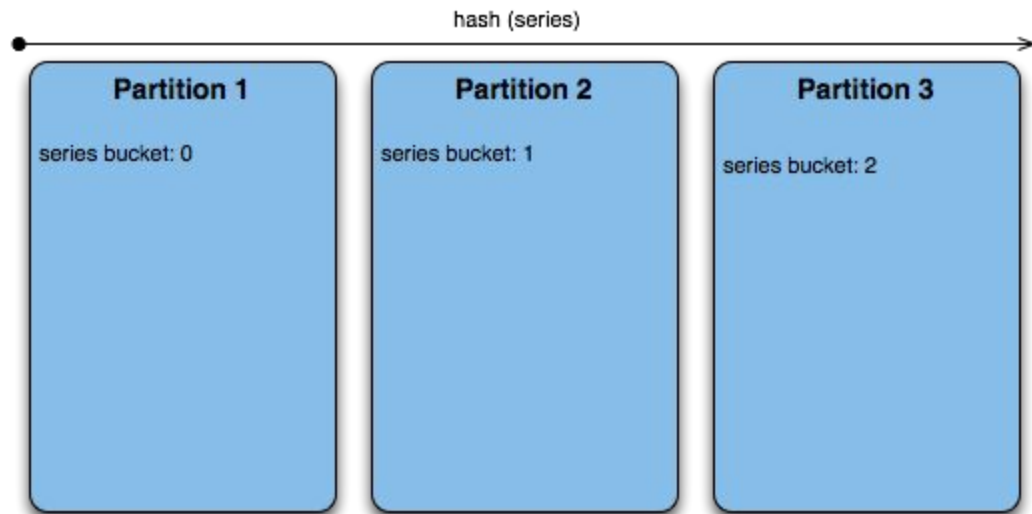


## Partitioning – By Series Hash (scans)



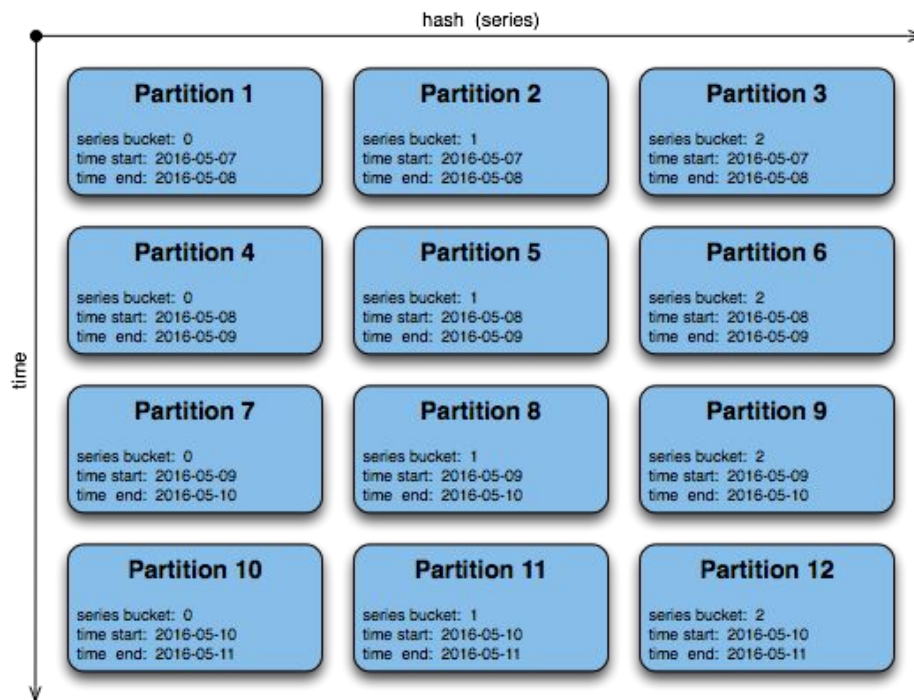
Scans are over a single partition

# Partitioning — By Series Hash



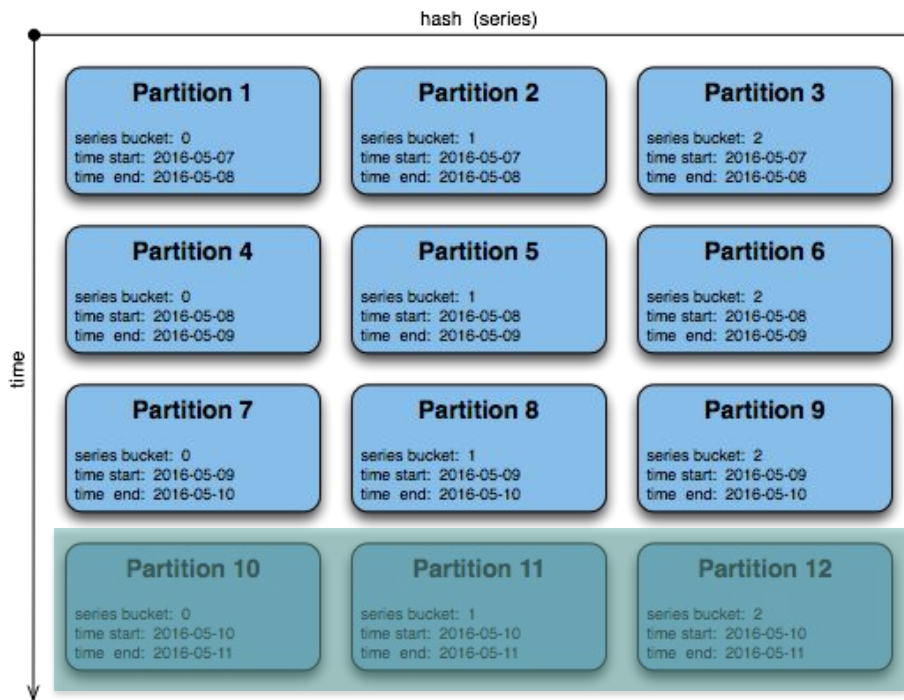
Partitions grow overtime, eventually becoming too big for a single server

# Partitioning — By Series Hash + Time Range



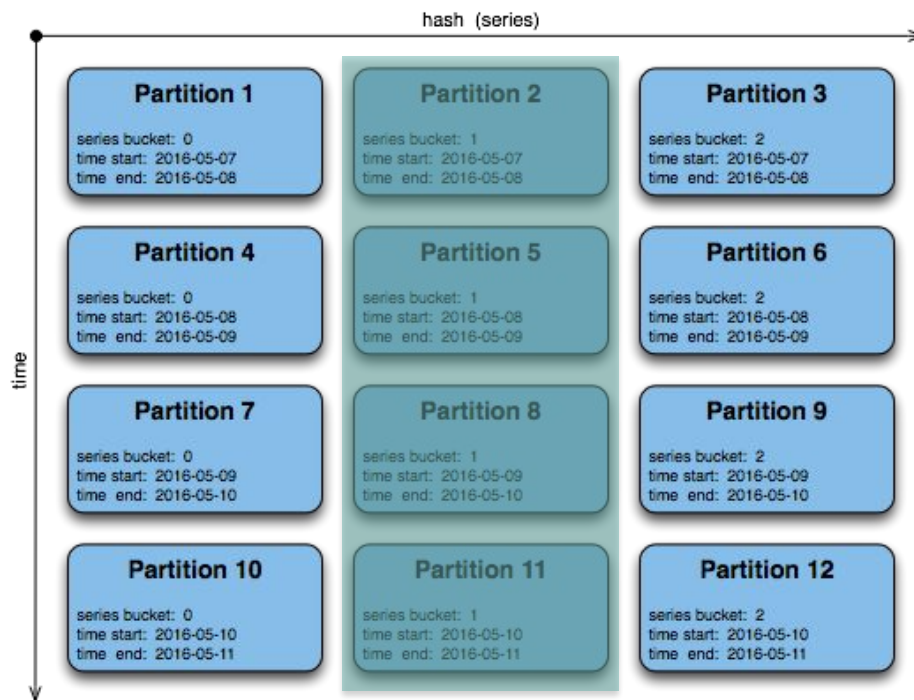


# Partitioning — By Series Hash + Time Range (inserts)



Inserts are spread among all partitions  
in the latest time range

# Partitioning — By Series Hash + Time Range (scans)



Big scans (across large time intervals)  
can be parallelized across partitions

---

# Spark SQL and Nifi walkthrough

- Deploy some Spark
- Deploy some Nifi
- ???
- Profit

---

# Spark DataSource optimizations

- Column projection and predicate pushdown
  - Only read the referenced columns
  - Convert 'WHERE' clauses into Kudu predicates
  - Kudu predicates automatically convert to primary key scans, etc

# Spark DataSource optimizations

## Predicate pushdown

```
scala> sqlContext.sql("select avg(value) from metrics where host = 'e1103.halxg.cloudera.com']").explain
== Physical Plan ==
TungstenAggregate(key=[], functions=[(avg(value#3),mode=Final,isDistinct=false)], output=[_c0#94])
+- TungstenExchange SinglePartition, None
   +- TungstenAggregate(key=[], functions=[(avg(value#3),mode=Partial,isDistinct=false)],
      output=[sum#98,count#99L])
      +- Project [value#3]
         +- Scan org.apache.kudu.spark.kudu.KuduRelation@e13cc49[value#3]
            PushedFilters: [EqualTo(host,e1103.halxg.cloudera.com)]
```

---

# Spark DataSource optimizations

## Partition pruning

```
scala> df.where("host like 'foo%'").rdd.partitions.length  
res1: Int = 20  
scala> df.where("host = 'foo'").rdd.partitions.length  
res2: Int = 1
```

---

# Writing via Spark

```
// Use KuduContext to create, delete, or write to Kudu tables
val kuduContext = new KuduContext("kudu-master:7051,kudu-master:7151,kudu-master:7251")

// Create a new Kudu table from a dataframe schema
// NB: No rows from the dataframe are inserted into the table
kuduContext.createTable("test_table", df.schema, Seq("key"), new CreateTableOptions().setNumReplicas(1))

// Insert, delete, upsert, or update data
kuduContext.insertRows(df, "test_table")
kuduContext.deleteRows(sqlContext.sql("select id from kudu_table where id >= 5"), "kudu_table")
kuduContext.upsertRows(df, "test_table")
kuduContext.updateRows(df.select("id", $"count" + 1, "test_table"))
```

---

# Spark SQL and Nifi walkthrough

<https://bit.ly/2ZsR6m6>

- Deploy some Spark
- Deploy some Nifi
- ???
- Profit



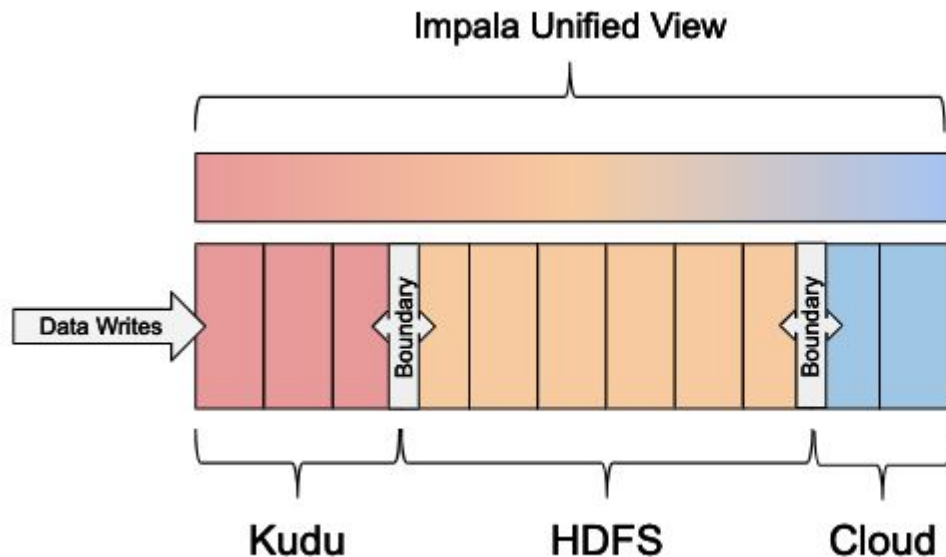
# Interesting pattern: Hierarchical storage

Return of the Lambda? (but longer time scale)

- Use Impala to periodically (e.g. every month) move data from Kudu into cold storage
- Query both hot Kudu data and colder HDFS and Cloud data with a view
- Simpler primitives than before

Excellent blog post:

<https://kudu.apache.org/2019/03/05/transparent-hierarchical-storage-management-with-apache-kudu-and-impala.html>



---

# Thank you!

Twitter: @ApacheKudu

Slack: <https://getkudu-slack.herokuapp.com/>

Website: [kudu.apache.org](https://kudu.apache.org)

Questions?

[awong@cloudera.com](mailto:awong@cloudera.com)